

Run-off triangles: Deterministic vs Stochastic methodologies

Davide Zicca

A digression about the Non-Life industry

In the Non-Life industry, the typical contracts expire in 12 months. But the claims payment process can take years or even decades. Moreover, often not even the delivery date of their product is known to insurers. The biggest item on the liabilities side of an insurer's balance sheet is often the provision or reserves for future claims payments. Those reserves can be broken down into case reserves (or outstanding claims), which are losses already reported to the insurance company and losses that are incurred but not reported (IBNR) yet. We have basically two different ways to assess the reserving problems: via deterministic or stochastic models. Solvency II in Europe, have fostered further research and promoted the use of stochastic and statistical techniques. In particular, for many countries, extreme percentiles of reserve deterioration over a fixed time period have to be estimated for the purpose of capital setting. I'll show how the data are structured for a typical reserving case. Then there's the discussion of the classical deterministic chain-ladder reserving method with an introduction of the concept of a tail factor. The chain-ladder algorithm can be considered a weighted linear regression through the origin, and move on from there to introduce stochastic reserving models. The Mack model, which provides a stochastic framework for the chain-ladder methods and allows the estimation of the mean squared error of the payment predictions. To estimate the full distribution of the reserve, we consider a bootstrap approach. We will do this analysis for two triangles taken from: https://www.casact.org/research/index.cfm?fa=loss_reserves_data and in particular, the csv file "PP Auto Data Set". I've then selected info that corresponds to "Germania Insurance Group" and "Battle Creek Mutual Insurance Co.". For which regards "Battle Creek Mutual Insurance Co.", I'll comment the results obtained, due to the fact that I've followed the same methodology applied to "Germania Insurance Group" (all infos are available in the R files and I've also saved all the plots that can be provided, if needed). We initialize our data, in order to have the first column that holds the origin year, the second column the development year and the third column has the incremental payments / transactions, for "Germania Insurance Group" in this way:

```
#####  
#                               Development Triangles                               #  
#####  
library(ChainLadder)  
  
##  
## Welcome to ChainLadder version 0.2.12  
##  
## To cite package 'ChainLadder' in publications use:  
##  
## Markus Gesmann, Daniel Murphy, Yanwei (Wayne) Zhang, Alessandro  
## Carrato, Mario Wuthrich, Fabio Concina and Eric Dal Moro (2021).  
## ChainLadder: Statistical Methods and Models for Claims Reserving in  
## General Insurance. R package version 0.2.12.  
## https://CRAN.R-project.org/package=ChainLadder  
##  
## To suppress this message use:  
## suppressPackageStartupMessages(library(ChainLadder))
```

```
library(lattice)
library(AER)
```

```
## Loading required package: car
## Loading required package: carData
## Loading required package: lmtest
## Loading required package: zoo
##
## Attaching package: 'zoo'
## The following objects are masked from 'package:base':
##
##      as.Date, as.Date.numeric
## Loading required package: sandwich
## Loading required package: survival
library(fitdistrplus)
```

```
## Loading required package: MASS
n<-10
Claims <- data.frame(originf = factor(rep(1988:1997, n:1)),
                      dev=sequence(n:1),
                      inc.paid=
                        c(2510,2653,2771,2799,2828,2810,2806,2803,2800,2804,
                          2500,2813,2903,2846,2839,2840,2871,2872,2873,
                          3648,4057,4059,4061,3996,3989,3983,3984,
                          4623,4998,4927,4831,4812,4804,4804,
                          5834,5585,5371,5216,5259,5242,
                          7248,6837,6698,6641,6564,
                          8757,8789,8542,8884,
                          9855,9234,9311,
                          10435,9604,
                          13768))

(inc.triangle <- with(Claims, {
  M <- matrix(nrow=n, ncol=n,
              dimnames=list(origin=levels(originf), dev=1:n))
  M[cbind(originf, dev)] <- inc.paid
  M
}))
```

```
##      dev
## origin  1    2    3    4    5    6    7    8    9   10
## 1988  2510 2653 2771 2799 2828 2810 2806 2803 2800 2804
## 1989  2500 2813 2903 2846 2839 2840 2871 2872 2873   NA
## 1990  3648 4057 4059 4061 3996 3989 3983 3984   NA   NA
## 1991  4623 4998 4927 4831 4812 4804 4804   NA   NA   NA
## 1992  5834 5585 5371 5216 5259 5242   NA   NA   NA   NA
## 1993  7248 6837 6698 6641 6564   NA   NA   NA   NA   NA
## 1994  8757 8789 8542 8884   NA   NA   NA   NA   NA   NA
## 1995  9855 9234 9311   NA   NA   NA   NA   NA   NA   NA
## 1996 10435 9604   NA   NA   NA   NA   NA   NA   NA   NA
```

```
## 1997 13768 NA NA NA NA NA NA NA NA NA
```

We will forecast the future claims development in the bottom right corner of the triangle and potential further developments beyond development age 10. Often it is helpful to consider the cumulative development of claims as well, which is presented below:

```
(cum.triangle <- t(apply(inc.triangle, 1, cumsum)))
```

```
##      dev
## origin 1 2 3 4 5 6 7 8 9 10
## 1988 2510 5163 7934 10733 13561 16371 19177 21980 24780 27584
## 1989 2500 5313 8216 11062 13901 16741 19612 22484 25357 NA
## 1990 3648 7705 11764 15825 19821 23810 27793 31777 NA NA
## 1991 4623 9621 14548 19379 24191 28995 33799 NA NA NA
## 1992 5834 11419 16790 22006 27265 32507 NA NA NA NA
## 1993 7248 14085 20783 27424 33988 NA NA NA NA NA
## 1994 8757 17546 26088 34972 NA NA NA NA NA NA
## 1995 9855 19089 28400 NA NA NA NA NA NA NA
## 1996 10435 20039 NA NA NA NA NA NA NA NA NA
## 1997 13768 NA NA NA NA NA NA NA NA NA NA
```

The latest diagonal of the triangle presents the latest cumulative paid position of all origin years:

```
(latest.paid <- cum.triangle[row(cum.triangle) == n - col(cum.triangle) + 1])
```

```
## [1] 13768 20039 28400 34972 33988 32507 33799 31777 25357 27584
```

We add the cumulative paid data as a column to the data frame as well:

```
Claims$cum.paid <- cum.triangle[with(Claims, cbind(originf, dev))]
```


To start the reserving analysis, we plot the data:

```
op <- par(fig=c(0,0.5,0,1), cex=0.8, oma=c(0,0,0,0))

with(Claims, {
  interaction.plot(x.factor=dev, trace.factor=originf, response=inc.paid,
    fun=sum, type="b", bty="n", legend=FALSE); axis(1, at=1:n)
  par(fig=c(0.45,1,0,1), new=TRUE, cex=0.8, oma=c(0,0,0,0))
  interaction.plot(x.factor=dev, trace.factor=originf, response=cum.paid,
    fun=sum, type="b", bty="n"); axis(1, at=1:n)
})

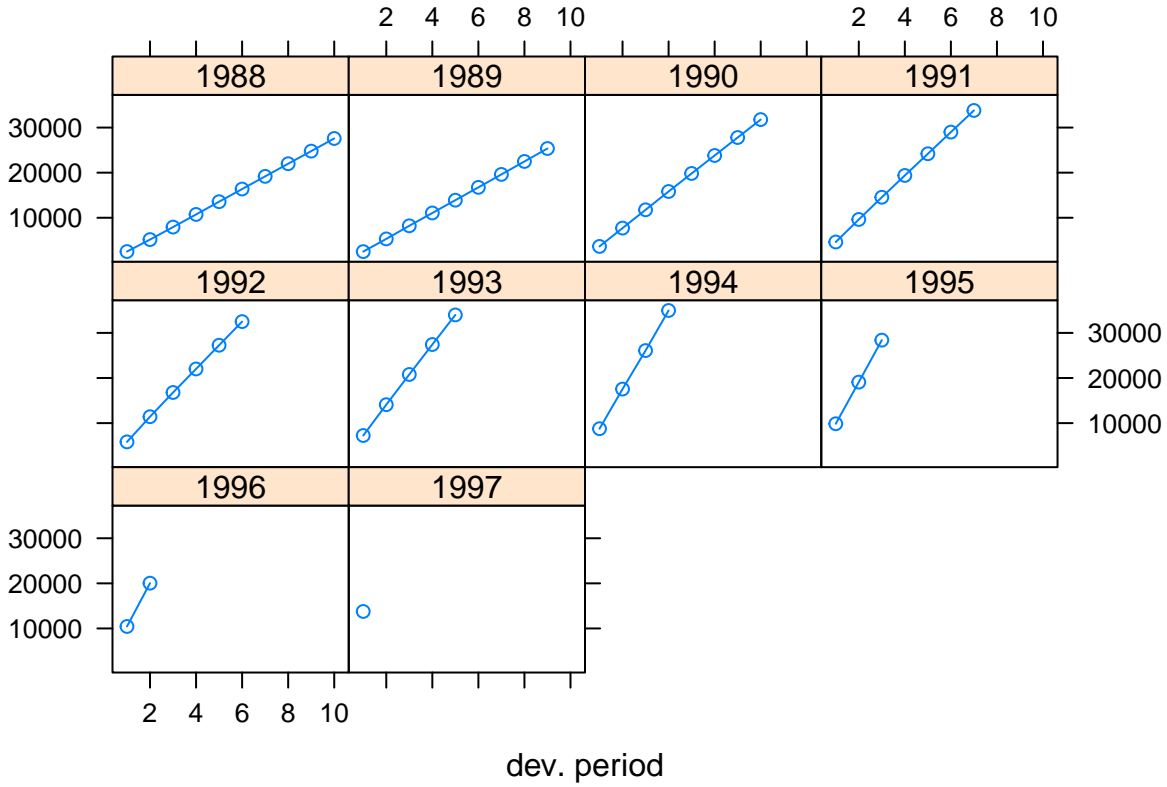
mtext("Incremental and cumulative claims development",
  side=3, outer=TRUE, line=-3, cex = 1.1, font=2)
```

Incremental and cumulative claims development



```
par(op)

library(lattice)
plot(as.triangle(cum.triangle),lattice=TRUE)
```



```
#xyplot(cum.paid ~ dev | originf, data=Claims, t="b", layout=c(4,2),
#as.table=TRUE, main="Cumulative claims development")
```

We can see the incremental and cumulative claims development by origin year. The triangle appears to be fairly well behaved. The last year, 1997, appear to be considerably higher than years 1994 to 1996. The second graph is about the cumulative claims developments by origin year using the lattice package, with one panel per origin year. Using terminology in Wuthrich & Merz (2008), cumulative payments are noted as

$$C_{i,j}$$

, for origin period i (or period of occurrence) seen after j periods of development and incremental payments

$$X_{i,j}$$

. The outstanding liabilities, or reserves, for accident year i at time j is given by:

$$R_{i,j} = \sum_{k>j} X_{i,k} = (\lim_{k \rightarrow \infty} C_{i,k}) - C_{i,j}$$

Because this quantity involves unobserved data (i.e. amounts that will be paid in the future),

$$R_{i,j}$$

will be the estimated claims reserves. Note that:

- For convenience, we will assume that we work on square matrices, with n rows and n columns.
- Throughout this chapter we note the first development period as 1. Other authors use 0 for the first development period. This assumption does not have any practical implications.

- Many of the methods and models presented here can be applied to paid and reported (often also called ‘incurred’) data. We either have to estimate the reserve or incurred but not reported (IBNR) claims. For the purpose of this chapter, we assume:

$$Ultimate\ loss\ cost = paid + reserve = paid + case\ reserve + IBNR = incurred + IBNR$$

- Some methods and models require data > 0 , which for paid claims should be given (insurers rarely receive money back from the insured after a claim was paid, such as late salvage or subrogation payments), but case reserves can show negative adjustments over time; therefore incremental incurred triangles do show negatives occasionally. To ensure that data have only positive values, it can be temporarily shifted, or, in a given context, be ignored.

As previously stated, I’ll now show the classical deterministic chain-ladder reserving method, also known as loss development factor (LDF) method. The idea is that, using this deterministic algorithm, we forecast claims based on historical data. It assumes that the proportional developments of claims from one development period to the next is the same for all origin periods. Most commonly as a first step, the age-to-age link ratios are calculated as the volume weighted average development ratios of a cumulative loss development triangle from one development period to the next

$$C_{i,k}, i, k = 1, \dots, n$$

$$f_k = \frac{\sum_{i=1}^{n-k} C_{i,k+1}}{\sum_{i=1}^{n-k} C_{i,k}}$$

```
f <- sapply((n-1):1, function(i) {
  sum( cum.triangle[1:i, n-i+1] ) / sum( cum.triangle[1:i, n-i] )
})
f
```

```
## [1] 1.984840 1.495681 1.332426 1.247094 1.199364 1.168349 1.145069 1.127586
## [9] 1.113156
```

Initially we expect no further development after year 10. Hence, we set the last link ratio (often called the tail factor) to 1:

```
tail <- 1

(f <- c(f, tail))
```

```
## [1] 1.984840 1.495681 1.332426 1.247094 1.199364 1.168349 1.145069 1.127586
## [9] 1.113156 1.000000
```

These factors f_k are then applied to the latest cumulative payment in each row

$$C_{i,n-i+1}$$

to produce stepwise forecasts for future payment years

$$k \in n - i + 1, \dots, n$$

:

$$\hat{C}_{i,k+1} = f_k \hat{C}_{i,k}$$

starting with

$$C_{i,n-i+1} = C_{i,n-i+1}$$

. The squaring of the claims triangle is calculated below:

```
full.triangle <- cum.triangle
```

```
for(k in 1:(n-1)){
  full.triangle[(n-k+1):n, k+1] <- full.triangle[(n-k+1):n,k]*f[k]
}
```

```
full.triangle
```

```
##      dev
## origin      1      2      3      4      5      6      7      8
## 1988 2510 5163.00 7934.00 10733.00 13561.00 16371.00 19177.00 21980.00
## 1989 2500 5313.00 8216.00 11062.00 13901.00 16741.00 19612.00 22484.00
## 1990 3648 7705.00 11764.00 15825.00 19821.00 23810.00 27793.00 31777.00
## 1991 4623 9621.00 14548.00 19379.00 24191.00 28995.00 33799.00 38702.20
## 1992 5834 11419.00 16790.00 22006.00 27265.00 32507.00 37979.51 43489.16
## 1993 7248 14085.00 20783.00 27424.00 33988.00 40763.98 47626.54 54535.68
## 1994 8757 17546.00 26088.00 34972.00 43613.38 52308.32 61114.35 69980.16
## 1995 9855 19089.00 28400.00 37840.89 47191.15 56599.37 66127.79 75720.90
## 1996 10435 20039.00 29971.94 39935.38 49803.19 59732.15 69787.97 79912.06
## 1997 13768 27327.28 40872.88 54460.07 67916.85 81457.02 95170.19 108976.46
##      dev
## origin      9      10
## 1988 24780.00 27584.00
## 1989 25357.00 28226.29
## 1990 35831.31 39885.83
## 1991 43640.07 48578.19
## 1992 49037.79 54586.70
## 1993 61493.69 68452.06
## 1994 78908.67 87837.65
## 1995 85381.85 95043.30
## 1996 90107.74 100303.95
## 1997 122880.36 136784.99
```

The last column contains the forecast ultimate loss cost:

```
(ultimate.paid <- full.triangle[,n])
```

```
##      1988      1989      1990      1991      1992      1993      1994      1995
## 27584.00 28226.29 39885.83 48578.19 54586.70 68452.06 87837.65 95043.30
##      1996      1997
## 100303.95 136784.99
```

The cumulative products of the age-to-age development ratios provide the loss development factors for the latest cumulative paid claims for each row to ultimate:

```
(ldf <- rev(cumprod(rev(f))))
```

```
## [1] 9.934993 5.005437 3.346595 2.511656 2.014007 1.679229 1.437267 1.255179
## [9] 1.113156 1.000000
```

The inverse of the loss development factor estimates the proportion of claims developed to date for each origin year, often also called the gross up factors or growth curve:

```
(dev.pattern <- 1/ldf)
```

```
## [1] 0.1006543 0.1997828 0.2988112 0.3981436 0.4965227 0.5955114 0.6957649
## [8] 0.7966989 0.8983469 1.0000000
```

The total estimated outstanding loss reserve with this method is:

```
(reserve <- sum (latest.paid * (ldf - 1)))
```

```
## [1] 405092
```

Note that: The basic chain-ladder algorithm has the implicit assumption that each origin period has its own unique level and that development factors are independent of the origin periods; or equivalently, there is a constant payment pattern. Therefore, if a_i is the ultimate (cumulative) claim for origin period i and b_j is the percentage of ultimate claims in development period j , with

$$\sum b_j = 1$$

, then the incremental payment

$$X_{i,j}$$

can be described as

$$X_{i,j} = a_i b_j$$

```
# otherwise
#sum(ultimate.paid - latest.paid)
```

```
a <- ultimate.paid
```

```
(b <- c(dev.pattern[1], diff(dev.pattern)))
```

```
## [1] 0.10065432 0.09912843 0.09902841 0.09933248 0.09837903 0.09898874
## [7] 0.10025346 0.10093408 0.10164792 0.10165313
```

```
(X.hat <- a %*% t(b))
```

```
##           [,1]      [,2]      [,3]      [,4]      [,5]      [,6]      [,7]
## [1,] 2776.449 2734.359 2731.600 2739.987 2713.687 2730.505 2765.392
## [2,] 2841.098 2798.028 2795.205 2803.787 2776.875 2794.085 2829.783
## [3,] 4014.681 3953.820 3949.831 3961.958 3923.929 3948.248 3998.693
## [4,] 4889.605 4815.480 4810.621 4825.392 4779.075 4808.694 4870.132
## [5,] 5494.387 5411.093 5405.634 5422.232 5370.186 5403.468 5472.505
## [6,] 6889.996 6785.545 6778.699 6799.513 6734.247 6775.983 6862.556
## [7,] 8841.239 8707.208 8698.423 8725.131 8641.382 8694.937 8806.028
## [8,] 9566.519 9421.493 9411.988 9440.887 9350.268 9408.216 9528.420
## [9,] 10096.026 9942.974 9932.942 9963.440 9867.805 9928.962 10055.819
## [10,] 13768.000 13559.281 13545.600 13587.191 13456.774 13540.173 13713.169
##           [,8]      [,9]      [,10]
## [1,] 2784.166 2803.856 2804.000
## [2,] 2848.995 2869.144 2869.291
## [3,] 4025.840 4054.312 4054.520
## [4,] 4903.195 4937.872 4938.125
## [5,] 5509.658 5548.624 5548.909
## [6,] 6909.146 6958.010 6958.366
## [7,] 8865.812 8928.514 8928.972
## [8,] 9593.108 9660.954 9661.449
## [9,] 10124.087 10195.688 10196.211
## [10,] 13806.267 13903.909 13904.622
```


Bornhuetter & Ferguson methodology

As the chain-ladder method is a deterministic algorithm and does not regard the observations as realizations of random variables but absolute values, the forecast of the most recent origin periods can be quite unstable. To address this issue, Bornhuetter & Ferguson (BF, 1972) suggested a credibility approach, which combines the chain-ladder forecast with prior information on expected loss costs, for example, from pricing data. Under this approach the chain-ladder development to ultimate pattern is used as weighting factors between the pure chain-ladder and expected loss cost estimates. Suppose the expected loss cost for the 1997 origin year is 20,000; then the BF method would estimate the ultimate loss cost as:

```
#Suppose the expected loss cost for the 1997 origin year is 20,000
(BF1997 <- ultimate.paid[n] * dev.pattern[1] + 20000 * (1 - dev.pattern[1]))

##      1997
## 31754.91
```

Tail Factors

In the previous section we implicitly assumed that there are no claims payment after 10 years, or in other words, that the oldest origin year is fully developed. However, often it is not suitable to assume that the oldest origin year is fully settled. A typical approach to overcome this shortcoming is to extrapolate the development ratios, for example, assuming a linear model of the log development ratios minus one, which reflects the incremental changes on the previous cumulative payments:

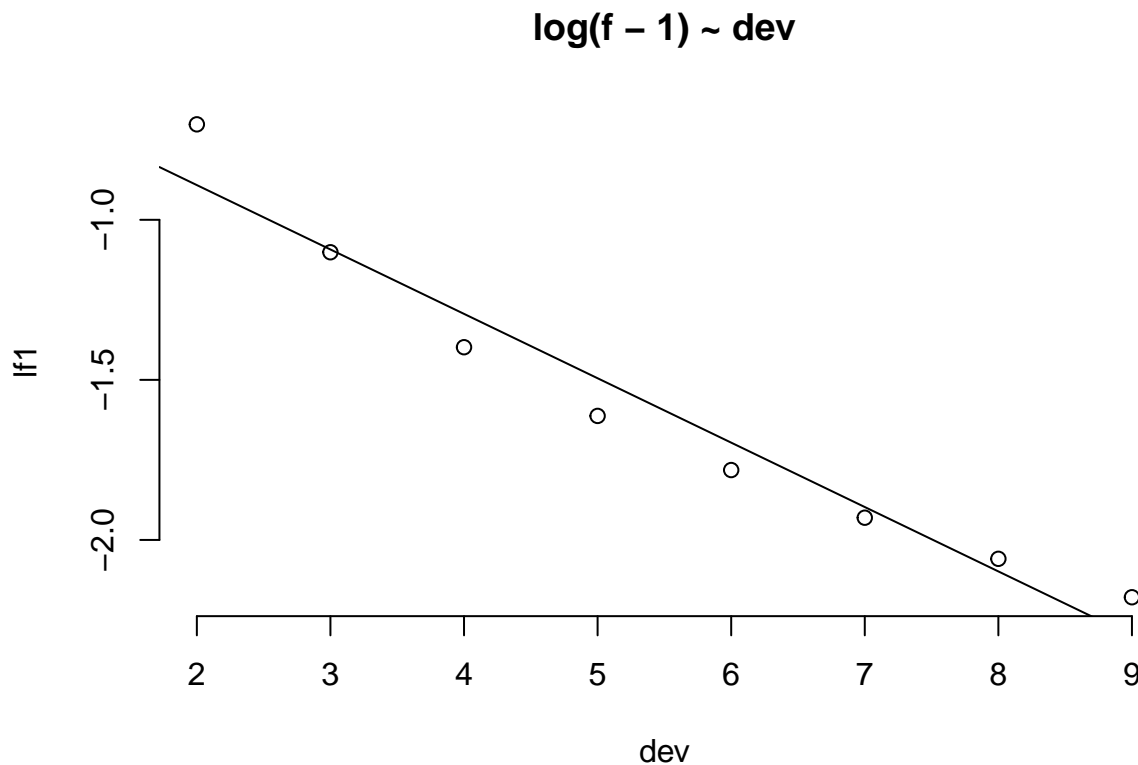
```
dat <- data.frame(lf1=log(f[-c(1,n)]-1), dev=2:(n-1))

(m <- lm(lf1 ~ dev , data=dat))

##
## Call:
## lm(formula = lf1 ~ dev, data = dat)
##
## Coefficients:
## (Intercept)      dev
##   -0.4893      -0.2011

plot(lf1 ~ dev, main="log(f - 1) ~ dev", data=dat, bty="n")

abline(m)
```



```
sigma <- summary(m)$sigma
extrapolation <- predict(m, data.frame(dev=n:100))
(tail <- prod(exp(extrapolation + 0.5*sigma^2) + 1))
```

```
## [1] 1.558258
```

#This is the plot of the loss development factors -1 on a log-scale against development period

More generally, the factors used to project the future payments need not always be drawn from the dollar weighted averages of the triangle. Other sources of factors from which the actuary may select link ratios include simple averages from the triangle, averages weighted toward more recent observations or adjusted for outliers, and benchmark patterns based on related, more credible loss experience. Also, because the ultimate value of claims is simply the product of the most current diagonal and the cumulative product of the link ratios, the completion of the interior of the triangle is usually not displayed; instead, the eventual value of the claims, or ultimate value is shown. For example, suppose the actuary decides that the volume weighted factors from the claims triangle are representative of expected future growth, but discards the tail factor derived from the linear fit in favour of a tail based on data from a larger book of similar business. The LDF method is as follows:

```
library(ChainLadder)
ata(cum.triangle)
```

```
##      dev
## origin 1-2  2-3  3-4  4-5  5-6  6-7  7-8  8-9  9-10
##   1988 2.057 1.537 1.353 1.263 1.207 1.171 1.146 1.127 1.113
```

##	1989	2.125	1.546	1.346	1.257	1.204	1.171	1.146	1.128	NA
##	1990	2.112	1.527	1.345	1.253	1.201	1.167	1.143	NA	NA
##	1991	2.081	1.512	1.332	1.248	1.199	1.166	NA	NA	NA
##	1992	1.957	1.470	1.311	1.239	1.192	NA	NA	NA	NA
##	1993	1.943	1.476	1.320	1.239	NA	NA	NA	NA	NA
##	1994	2.004	1.487	1.341	NA	NA	NA	NA	NA	NA
##	1995	1.937	1.488	NA	NA	NA	NA	NA	NA	NA
##	1996	1.920	NA	NA	NA	NA	NA	NA	NA	NA
##	smpl	2.015	1.505	1.335	1.250	1.201	1.169	1.145	1.128	1.113
##	vwtd	1.985	1.496	1.332	1.247	1.199	1.168	1.145	1.128	1.113

Stochastic Reserving Models

As the provision for outstanding claims is often the biggest item on the liabilities side of an insurer's balance sheet, it is important not only to estimate the mean but also the uncertainty of the reserve. The key idea is to regard the observed data as one realization of a random variable, rather than absolutes. Statistical techniques also allow for more formal testing, make modelling assumptions more explicit and, in particular, help to monitor actual versus expected claims developments (A versus E). It is the regular A versus E exercise which can help to drive management actions. Hence, as a minimum, not only the mean reserve, or best estimate liabilities³, should be estimated but also the volatility of reserves. In this section we first show that the deterministic chain-ladder algorithm of the previous section can be considered a weighted linear regression through the origin. Indeed, the following Mack model provides a stochastic framework for the chain-ladder method and allows us to estimate the mean squared error of future payments, using many estimators from the linear regression output. An alternative to the Mack model is the Poisson model, a generalized linear model that replicates the chain-ladder forecasts as well. Yet, the Poisson model is often not directly applicable to insurance data, as the variance of the data is frequently greater than the mean and hence we consider a quasi-Poisson model to estimate uncertainty metrics. Following this, we present a bootstrap technique to estimate the full reserve distribution.

Chain-Ladder in the Context of Linear Regression Barnett & Zehnwirth (2000)

Let $C_{\cdot,k}$ denote the k-th column in the cumulative claims triangle. The chain-ladder algorithm can be seen as:

$$C_{\cdot,k+1} = f_k C_{\cdot,k} + \epsilon(k) \sim \mathcal{N}(0, \sigma_k^2 C_{\cdot,k}^\delta)$$

The parameter f_k describes the slope or the 'best' line through the origin and data points

$$(C_{\cdot,k}, C_{\cdot,k+1})$$

, with δ as a 'weighting' parameter. Barnett & Zehnwirth (2000) distinguish the cases:

- $\delta=0$, ordinary regression with intercept 0
- $\delta=1$, historical chain ladder age-to-age link ratios
- $\delta=2$, straight averages of the individual link ratios

Indeed, we can demonstrate the different cases by applying different linear models to our data. First, we add columns to the original dataframe Claims, to have payments of the current and previous development period next to each other; additionally we add a column with the development period as a factor.

```
#####
#                               Stochastic Reserving Models                               #
#                               Chain-Ladder in the Context of Linear Regression           #
#####

names(Claims)[3:4] <- c("inc.paid.k", "cum.paid.k")

ids <- with(Claims, cbind(originf, dev))

Claims <- within(Claims,{
  cum.paid.kp1 <- cbind(cum.triangle[, -1], NA)[ids]
  inc.paid.kp1 <- cbind(inc.triangle[, -1], NA)[ids]
  devf <- factor(dev)
})
```

In the next step we apply the linear regression function `lm` to each development period, vary the weighting parameter `delta` from 0 to 2 and extract the slope coefficients:

```
delta <- 0:2

ATA <- sapply(delta, function(d)
  coef(lm(cum.paid.kp1 ~ 0 + cum.paid.k : devf,
    weights=1/cum.paid.k^d, data=Claims))
)

dimnames(ATA)[[2]] <- paste("Delta = ", delta)

ATA
```

```
##          Delta = 0 Delta = 1 Delta = 2
## cum.paid.k:devf1  1.965004  1.984840  2.015225
## cum.paid.k:devf2  1.489876  1.495681  1.505314
## cum.paid.k:devf3  1.331220  1.332426  1.335315
## cum.paid.k:devf4  1.244770  1.247094  1.249881
## cum.paid.k:devf5  1.198127  1.199364  1.200722
## cum.paid.k:devf6  1.167763  1.168349  1.168966
## cum.paid.k:devf7  1.144806  1.145069  1.145317
## cum.paid.k:devf8  1.127589  1.127586  1.127584
## cum.paid.k:devf9  1.113156  1.113156  1.113156
```

Indeed, the development ratios for

$$\delta = 1$$

and

$$\delta = 2$$

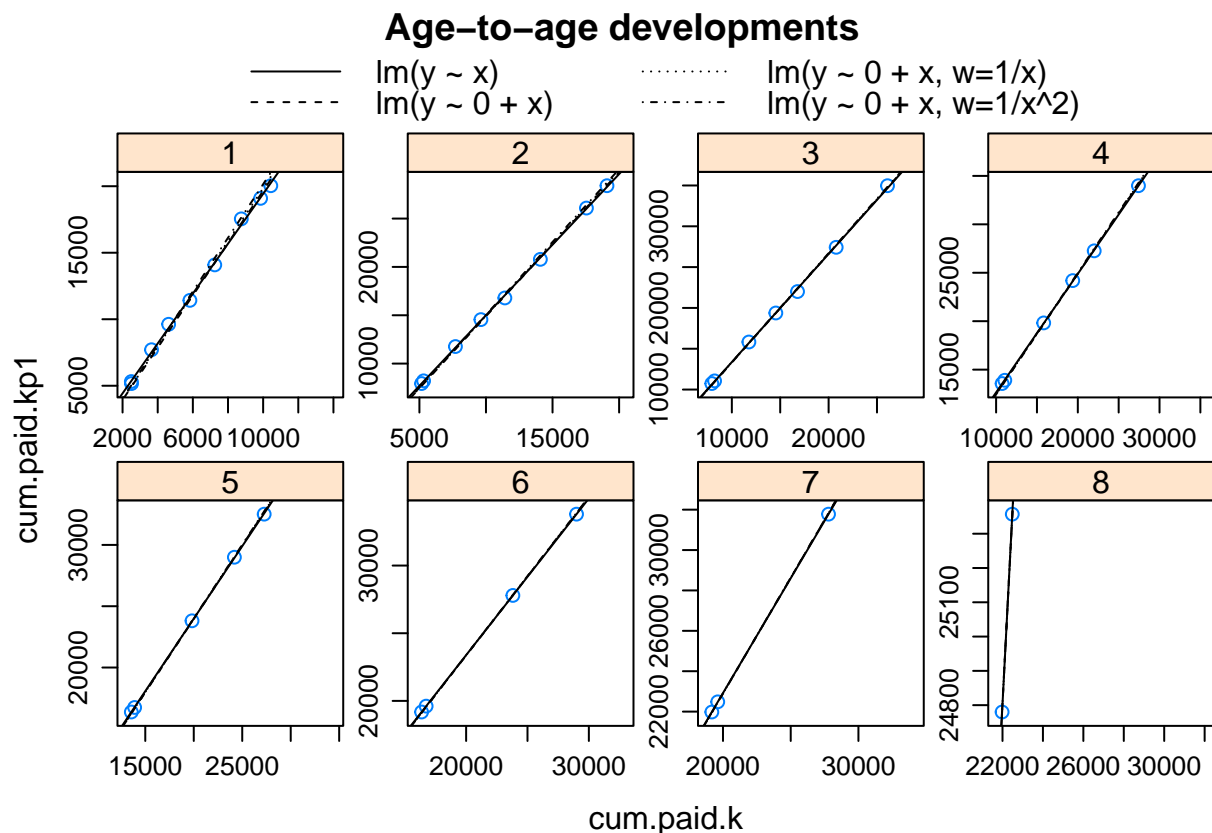
correspond with those of the previous section. The following graph is the plot of the cumulative development positions from one development year to the next for each development year, including regression lines of different linear models:

```
xyplot(cum.paid.kp1 ~ cum.paid.k | devf,
  data=subset(Claims, dev < (n-1)),
  main="Age-to-age developments", as.table=TRUE,
  scales=list(relation="free"),
  key=list(columns=2, lines=list(lty=1:4, type="l"),
    text=list(lab=c("lm(y ~ x)",
```

```

"lm(y ~ 0 + x)",
"lm(y ~ 0 + x, w=1/x)",
"lm(y ~ 0 + x, w=1/x^2)"))),
panel=function(x,y,...){
  panel.xyplot(x,y,...)
  if(length(x)>1){
    panel.abline(lm(y ~ x), lty=1)
    panel.abline(lm(y ~ 0 + x), lty=2)
    panel.abline(lm(y ~ 0 + x, weights=1/x), lty=3)
    panel.abline(lm(y ~ 0 + x, weights=1/x^2), lty=4)
  }
}
)

```



Note that for all development periods we observe no difference in the slope of the linear regression with and without an intercept.

Mack Model

Mack (1993, 1999) suggested a model to estimate the first two moments (mean and standard errors) of the chain-ladder forecast, without assuming a distribution under three conditions. Note that Mack uses the following notation for the weighting parameter $\alpha = 2 - \delta$, with δ defined as in the previous section. In other words, the Mack model assumes that the link ratios for each development period are consistent across all origin periods (CL1), the volatility decreases as losses are paid (CL2) and all origin periods are independent; for example, there is no structural change or market cycle (CL3). If these assumptions hold, the Mack model gives an unbiased estimator for future claims.

In R we have the following output:

```
#####
#                               Mack Model                               #
#####

library(ChainLadder)

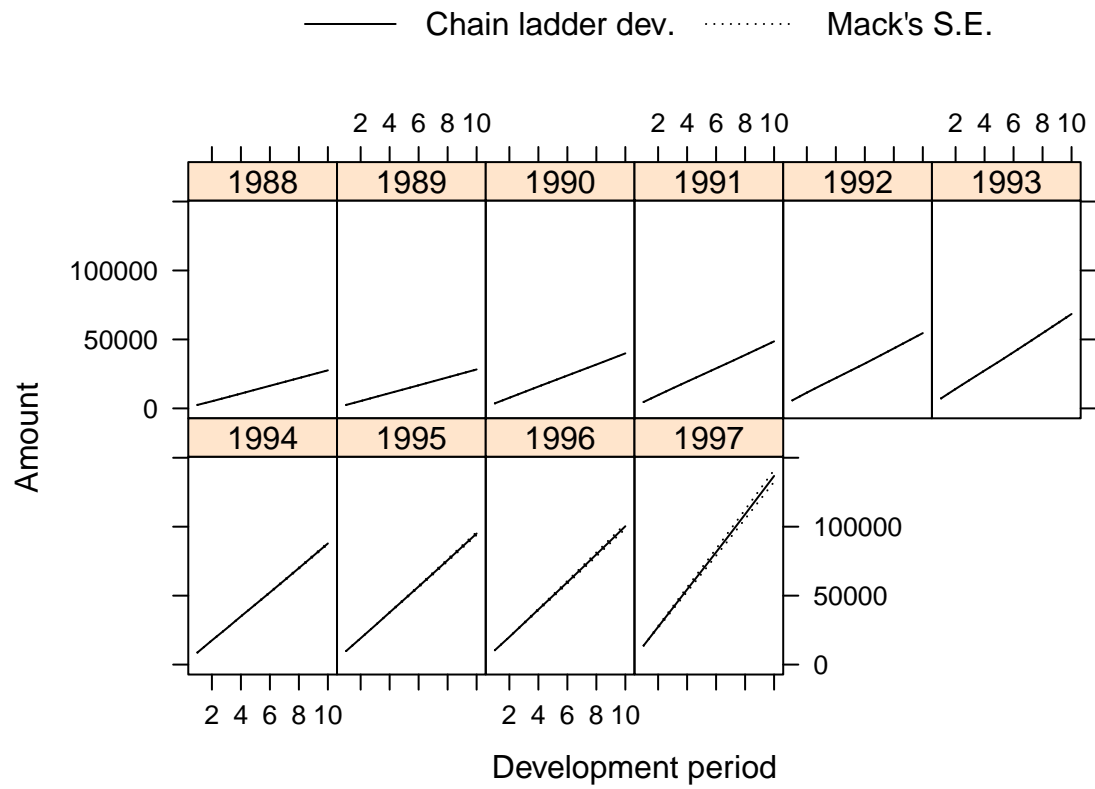
(mack <- MackChainLadder(cum.triangle, weights=1, alpha=1,est.sigma="Mack"))

## MackChainLadder(Triangle = cum.triangle, weights = 1, alpha = 1,
##   est.sigma = "Mack")
##
##      Latest Dev.To.Date Ultimate      IBNR Mack.S.E CV(IBNR)
## 1988 27,584      1.000  27,584         0      0.00      NaN
## 1989 25,357      0.898  28,226    2,869      1.44 0.000503
## 1990 31,777      0.797  39,886    8,109     10.88 0.001342
## 1991 33,799      0.696  48,578   14,779     76.66 0.005187
## 1992 32,507      0.596  54,587   22,080    154.56 0.007000
## 1993 33,988      0.497  68,452   34,464    342.22 0.009930
## 1994 34,972      0.398  87,838   52,866    666.46 0.012607
## 1995 28,400      0.299  95,043   66,643   1,116.37 0.016751
## 1996 20,039      0.200 100,304   80,265   1,793.19 0.022341
## 1997 13,768      0.101 136,785  123,017   4,265.46 0.034674
##
##              Totals
## Latest:    282,191.00
## Dev:              0.41
## Ultimate: 687,282.96
## IBNR:       405,091.96
## Mack.S.E    5,305.39
## CV(IBNR):      0.01
```

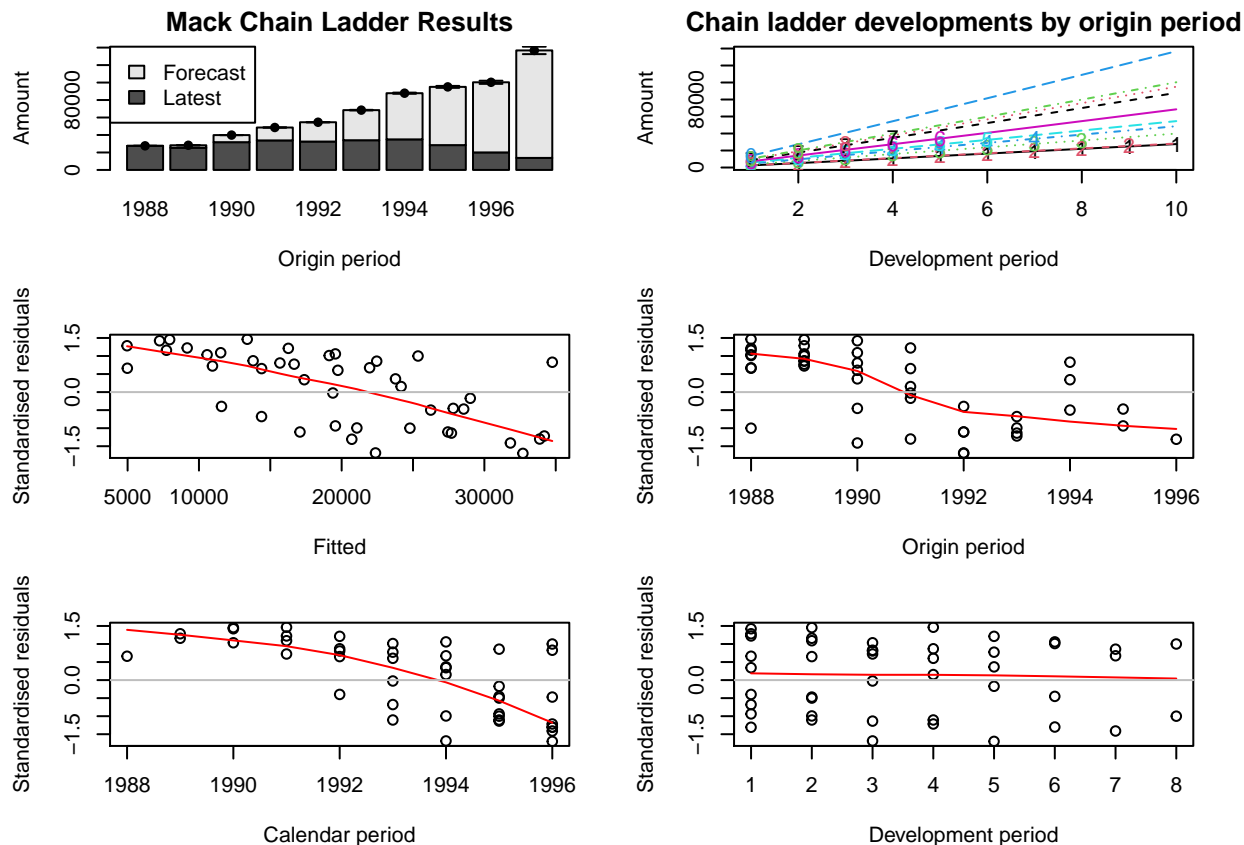
Plot of the actual and expected cumulative claims development and estimated standard error of the Mack model forecast:

```
plot(mack, lattice=TRUE, layout=c(6,2))
```

Chain ladder developments by origin period



```
plot(mack)
```



The top-left panel shows the latest actual position with the forecasts stacked on top and whiskers indicating the estimated standard error. The top-right panel presents the claims developments to ultimate for each origin year. The four residual plots show the standardized residuals against fitted values, origin, calendar and development period. The residual plots should not show any obvious patterns and about 95% of the standardized residuals should be contained in the range of -2 to 2 for the Mack model to be strictly applicable (and this is exactly what we got).

Bootstrap Chain-Ladder

An alternative to asymptotic econometric relationships can be to use the bootstrap methodology. We use a two-stage simulation approach, following P.D. England & R.J. Verrall (2002). In the first stage, a quasi-Poisson model is applied to the claims triangle to forecast future payments. From this we calculate the scaled-Pearson residuals, assuming that they are approximately independent and identical distributed. These residuals are re-sampled with replacement many times to generate bootstrapped (pseudo) triangles and to forecast future claims payments to estimate the parameter error. Recall that the predictions of the quasi-Poisson model are the same as those from the chain-ladder method, hence we use the latter faster algorithm. In the second stage, we simulate the process error with the bootstrap value as the mean and an assumed process distribution, here a quasi-Poisson. The set of reserves obtained in this way forms the predictive distribution, from which summary statistics such as mean, prediction error or quantiles can be derived.

This two-stage bootstrapping/simulation approach is implemented in the `BootChainLadder` function as part of the `ChainLadder` package. As input parameters we provide the cumulative triangle, the number of bootstraps and the process distribution to be assumed:

```
#####
#                               #
#                               #
#####
```



```
set.seed(1)
```

```
(B <- BootChainLadder(cum.triangle, R=1000, process.distr="od.pois"))
```

```
## BootChainLadder(Triangle = cum.triangle, R = 1000, process.distr = "od.pois")
```

```
##
```





##	Latest	Mean Ultimate	Mean IBNR	IBNR.S.E	IBNR 75%	IBNR 95%
## 1988	27,584	27,584	0	0	0	0
## 1989	25,357	28,227	2,870	217	3,011	3,216
## 1990	31,777	39,896	8,119	389	8,393	8,730
## 1991	33,799	48,569	14,770	551	15,170	15,648
## 1992	32,507	54,583	22,076	716	22,559	23,210
## 1993	33,988	68,466	34,478	956	35,082	36,033
## 1994	34,972	87,821	52,849	1,328	53,715	54,963
## 1995	28,400	95,061	66,661	1,685	67,860	69,398
## 1996	20,039	100,262	80,223	2,108	81,636	83,516
## 1997	13,768	136,686	122,918	3,710	125,483	129,038

```
##
```

##	Totals
## Latest:	282,191
## Mean Ultimate:	687,155
## Mean IBNR:	404,964
## IBNR.S.E	7,366
## Total IBNR 75%:	410,307
## Total IBNR 95%:	416,818

IBNR 75% is the quantile which is useful in some line of business due to capital requirement purposes. In the 75% of the cases you'll have that IBNR and so you must set aside that amount of value. The same meaning is for IBNR 95%.

```
plot(B)
```



```
quantile(B, c(0.75,0.95,0.99, 0.995))
```

```
## $ByOrigin
##      IBNR 75%  IBNR 95%  IBNR 99%  IBNR 99.5%
## 1988      0.00      0.00      0.00      0.00
## 1989    3011.00    3216.05    3358.11    3418.26
## 1990    8393.00    8730.15    8947.05    9035.02
## 1991   15170.25   15647.65   16037.24   16130.21
## 1992   22558.50   23210.15   23654.06   23871.65
## 1993   35082.25   36033.00   36803.07   37106.50
## 1994   53715.00   54962.50   56216.16   56396.76
## 1995   67859.75   69397.80   70657.85   70908.18
## 1996   81635.50   83516.05   85015.75   85355.81
## 1997  125482.50  129037.55  130787.56  131834.02
##
## $Totals
##      Totals
## IBNR 75%:  410307.0
## IBNR 95%:  416817.8
## IBNR 99%:  421419.0
## IBNR 99.5%: 423251.8
```

Losses may follow a lognormal distribution. We can test this idea for our data by fitting a log-normal distribution to the predicted future payments. The `fitdistrplus` package by Delignette-Muller et al. (2013) makes it a one line in R:

```
library(fitdistrplus)
```

```
(fit <- fitdistr(B$IBNR.Totals[B$IBNR.Totals>0], "lnorm"))
```

```
## Fitting of the distribution 'lnorm' by maximum likelihood
```

```
## Parameters:
```

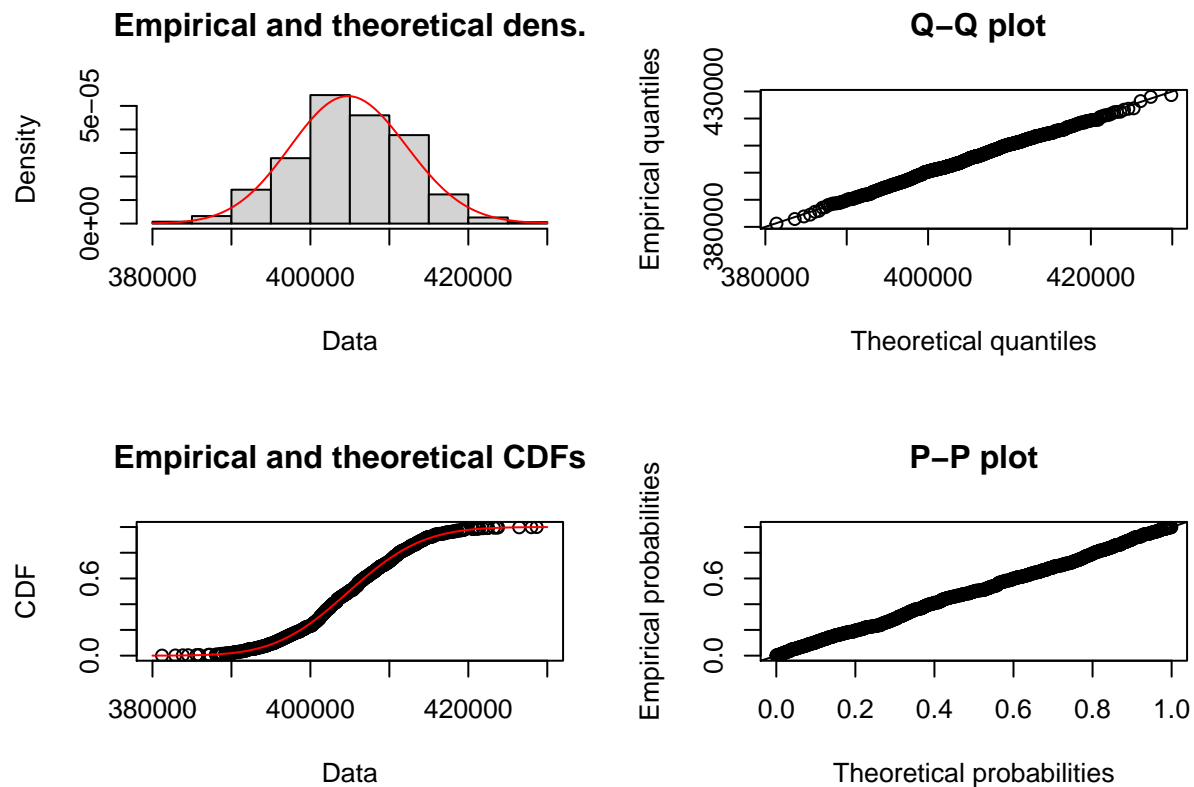
```
##           estimate   Std. Error
```

```
## meanlog 12.91138705 0.0005753809
```

```
## sdlog    0.01819514 0.0004013392
```

Plots of simulated data from BootChainLadder and a fitted log-normal distribution:

```
plot(fit)
```



```
qlnorm(0.995, fit$estimate["meanlog"], fit$estimate["sdlog"])
```

```
## [1] 424324.9
```

```
ny <- (col(inc.triangle) == (nrow(inc.triangle) - row(inc.triangle) + 2))
```

```
paid.ny <- apply(B$IBNR.Triangles, 3,
```

```
  function(x){
```

```
    next.year.paid <- x[col(x) == (nrow(x) - row(x) + 2)]
```

```
    sum(next.year.paid)
```

```
  })
```

```
paid.ny.995 <- B$IBNR.Triangles[, , order(paid.ny)[round(B$R*0.995)]]
```

```
inc.triangle.ny <- inc.triangle
```

```
(inc.triangle.ny[ny] <- paid.ny.995[ny])
```

```
## [1] 14776 10003 9485 9218 7156 5351 5126 4242 2873
```

The fit looks very reasonable.

Munich chain-ladder

It's a reserving method that reduces the gap between IBNR projections based on paid losses and IBNR projections based on incurred losses. The Munich chain-ladder method uses correlations between paid and incurred losses of the historical data into the projection for the future. I'll report the codes that I've run on R:

```
Incurred= cum.triangle
n<-10
Claimspaid <- data.frame(originf = factor(rep(1988:1997, n:1)),
                          dev=sequence(n:1),
                          cum.paid=
                            c(1586,2307,2604,2702,2745,2809,2806,2803,2800,2797,
                              1190,2044,2641,2756,2796,2818,2871,2872,2873,
                              2007,3281,3909,3955,3963,3982,3983,3984,
                              2300,4150,4631,4762,4782,4804,4804,
                              2685,4473,4982,5136,5203,5232,
                              3212,5409,6102,6439,6526,
                              3881,6833,7957,8357,
                              4804,7373,8486,
                              4664,7369,
                              5915))

(cumpaid.triangle <- with(Claimspaid, {
  M <- matrix(nrow=n, ncol=n,
              dimnames=list(origin=levels(originf), dev=1:n))
  M[cbind(originf, dev)] <- cum.paid
  M
}))
```

```
##      dev
## origin  1    2    3    4    5    6    7    8    9   10
## 1988 1586 2307 2604 2702 2745 2809 2806 2803 2800 2797
## 1989 1190 2044 2641 2756 2796 2818 2871 2872 2873   NA
## 1990 2007 3281 3909 3955 3963 3982 3983 3984   NA   NA
## 1991 2300 4150 4631 4762 4782 4804 4804   NA   NA   NA
## 1992 2685 4473 4982 5136 5203 5232   NA   NA   NA   NA
## 1993 3212 5409 6102 6439 6526   NA   NA   NA   NA   NA
## 1994 3881 6833 7957 8357   NA   NA   NA   NA   NA   NA
## 1995 4804 7373 8486   NA   NA   NA   NA   NA   NA   NA
## 1996 4664 7369   NA   NA   NA   NA   NA   NA   NA   NA
## 1997 5915   NA   NA   NA   NA   NA   NA   NA   NA   NA
```

```
(cum.paid.triangle <- t(apply(cumpaid.triangle, 1, cumsum)))
```

```
##      dev
## origin  1    2    3    4    5    6    7    8    9   10
## 1988 1586 3893 6497 9199 11944 14753 17559 20362 23162 25959
## 1989 1190 3234 5875 8631 11427 14245 17116 19988 22861   NA
## 1990 2007 5288 9197 13152 17115 21097 25080 29064   NA   NA
## 1991 2300 6450 11081 15843 20625 25429 30233   NA   NA   NA
## 1992 2685 7158 12140 17276 22479 27711   NA   NA   NA   NA
```

```
## 1993 3212 8621 14723 21162 27688 NA NA NA NA NA
## 1994 3881 10714 18671 27028 NA NA NA NA NA NA
## 1995 4804 12177 20663 NA NA NA NA NA NA NA
## 1996 4664 12033 NA NA NA NA NA NA NA NA NA
## 1997 5915 NA NA NA NA NA NA NA NA NA NA
```

```
Paid=cum.paid.triangle
```

```
MCL=MunichChainLadder(Paid=Paid ,Incurred= Incurred,
                      est.sigmaP = "log-linear", est.sigmaI = "log-linear",
                      tailP=FALSE, tailI=FALSE)
```

```
MCL
```

```
## MunichChainLadder(Paid = Paid, Incurred = Incurred, est.sigmaP = "log-linear",
##                   est.sigmaI = "log-linear", tailP = FALSE, tailI = FALSE)
```

```
##
```

```
## Latest Paid Latest Incurred Latest P/I Ratio Ult. Paid Ult. Incurred
## 1988 25,959 27,584 0.941 25,959 27,584
## 1989 22,861 25,357 0.902 25,640 28,224
## 1990 29,064 31,777 0.915 37,127 39,888
## 1991 30,233 33,799 0.894 44,918 48,574
## 1992 27,711 32,507 0.852 49,454 54,476
## 1993 27,688 33,988 0.815 61,380 68,131
## 1994 27,028 34,972 0.773 78,450 87,219
## 1995 20,663 28,400 0.728 85,713 94,567
## 1996 12,033 20,039 0.600 87,935 98,376
## 1997 5,915 13,768 0.430 117,421 132,024
```

```
## Ult. P/I Ratio
```

```
## 1988 0.941
## 1989 0.908
## 1990 0.931
## 1991 0.925
## 1992 0.908
## 1993 0.901
## 1994 0.899
## 1995 0.906
## 1996 0.894
## 1997 0.889
```


```
##
```

```
## Totals
```


```
## Paid Incurred P/I Ratio
## Latest: 229,155 282,191 0.81
## Ultimate: 613,997 679,064 0.90
```

```
plot(MCL)
```


Munich Chain Ladder Results




Munich Chain Ladder vs. Standard Chain L



Paid residual plot



Incurred residual plot



Recalling that Ultimate loss cost = paid + reserve = paid + case reserve + IBNR = incurred + IBNR:

- Mack ChainLadder -> Ultimate= 687,282.96
- Munich ChainLadder -> Ultimate= 679,064
- Bootstrap ChainLadder -> (Mean) Ultimate= 687,267

As a final comment between ChainLadder-based models, we can say that the Munich ChainLadder underestimated the value of the Ultimate loss cost w.r.to the other 2 models that I've taken into account. In terms of Reserving:

- Basic ChainLadder-> Reserve= 405092
- Mack ChainLadder -> Reserve= 405,091.96
- Bootstrap ChainLadder -> Reserve= 405,076
- Munich ChainLadder -> Reserve= 396,873